

# Vectorization

---

**Write-up 1:** Look at the assembly code above. The compiler has translated the code to set the start index at  $-2^{16}$  and adds to it for each memory access. Why doesn't it set the start index to 0 and use small positive offsets?

答：每一次循环，i都需要和SIZE（即2的16次方）作比较。先将内存置为-2的16次方，加上i后直接判断正负来比较大小，可以减少开销。

**Write-up 2:** This code is still not aligned when using AVX2 registers. Fix the code to make sure it uses aligned moves for the best performance.

答：修改代码如下，使程序使用32位对齐寄存器性能最优

```
#include <stdint.h>
#include <stdlib.h>
#include <math.h>

#define SIZE (1L << 32)

void test(uint8_t* restrict a, uint8_t* restrict b) {
    a = __builtin_assume_aligned(a, 32);
    b = __builtin_assume_aligned(b, 32);

    for (uint64_t i = 0; i < SIZE; i++) {
        a[i] += b[i];
    }
}
```

**Write-up 3:** Provide a theory for why the compiler is generating dramatically different assembly.

答：三元运算符不会按元素比较数组。

**Write-up 4:** Inspect the assembly and determine why the assembly does not include instructions with vector registers. Do you think it would be faster if it did vectorize? Explain.

答:

运行生成的汇编代码如下:

```
.text
.file    "example3.c"
.globl   test                                # -- Begin function test
.p2align 4, 0x90
.type    test,@function
test:                                         # @test
.Lfunc_begin0:
.file    1 "/home/zhangkeer/a4/recitation3" "example3.c"
.loc     1 9 0                               # example3.c:9:0
.cfi_startproc
# %bb.0:
#DEBUG_VALUE: test:a <- $rdi
#DEBUG_VALUE: test:a <- $rdi
#DEBUG_VALUE: test:b <- $rsi
#DEBUG_VALUE: test:b <- $rsi
pushq    %rax
.cfi_def_cfa_offset 16
.Ltmp0:
#DEBUG_VALUE: test:i <- 0
.loc     1 12 3 prologue_end                # example3.c:12:3
addq     $1, %rsi
.Ltmp1:
.loc     1 13 10                             # example3.c:13:10
movl     $65536, %edx                       # imm = 0x10000
#DEBUG_VALUE: test:a <- $rdi
callq    memcpy
.Ltmp2:
#DEBUG_VALUE: test:i <- undef
.loc     1 15 1                               # example3.c:15:1
popq     %rax
.cfi_def_cfa_offset 8
retq
.Ltmp3:
.Lfunc_end0:
.size    test, .Lfunc_end0-test
.cfi_endproc
# -- End function

.file    2 "/usr/include/x86_64-linux-gnu/bits" "types.h"
.file    3 "/usr/include/x86_64-linux-gnu/bits" "stdint-uintn.h"
.section .debug_str,"MS",@progbits,1
.Linfo_string0:
.asciz   "clang version 10.0.0-4ubuntu1 " # string offset=0
.Linfo_string1:
.asciz   "example3.c"                     # string offset=31
.Linfo_string2:
.asciz   "/home/zhangkeer/a4/recitation3" # string offset=42
.Linfo_string3:
.asciz   "test"                           # string offset=73
.Linfo_string4:
.asciz   "a"                             # string offset=78
.Linfo_string5:
.asciz   "unsigned char"                  # string offset=80
```

```

.Linfo_string6:
    .asciz  "__uint8_t"           # string offset=94
.Linfo_string7:
    .asciz  "uint8_t"           # string offset=104
.Linfo_string8:
    .asciz  "b"                 # string offset=112
.Linfo_string9:
    .asciz  "i"                 # string offset=114
.Linfo_string10:
    .asciz  "long unsigned int"  # string offset=116
.Linfo_string11:
    .asciz  "__uint64_t"        # string offset=134
.Linfo_string12:
    .asciz  "uint64_t"          # string offset=145
    .section    .debug_loc,"",@progbits
.Ldebug_loc0:
    .quad    .Lfunc_begin0-.Lfunc_begin0
    .quad    .Ltmp2-.Lfunc_begin0
    .short   1                  # Loc expr size
    .byte    85                 # DW_OP_reg5
    .quad    0
    .quad    0
.Ldebug_loc1:
    .quad    .Lfunc_begin0-.Lfunc_begin0
    .quad    .Ltmp1-.Lfunc_begin0
    .short   1                  # Loc expr size
    .byte    84                 # DW_OP_reg4
    .quad    0
    .quad    0
.Ldebug_loc2:
    .quad    .Ltmp0-.Lfunc_begin0
    .quad    .Ltmp2-.Lfunc_begin0
    .short   2                  # Loc expr size
    .byte    48                 # DW_OP_lit0
    .byte    159                # DW_OP_stack_value
    .quad    0
    .quad    0
    .section    .debug_abbrev,"",@progbits
    .byte    1                  # Abbreviation Code
    .byte    17                 # DW_TAG_compile_unit
    .byte    1                  # DW_CHILDREN_yes
    .byte    37                 # DW_AT_producer
    .byte    14                 # DW_FORM_strp
    .byte    19                 # DW_AT_language
    .byte    5                  # DW_FORM_data2
    .byte    3                  # DW_AT_name
    .byte    14                 # DW_FORM_strp
    .byte    16                 # DW_AT_stmt_list
    .byte    23                 # DW_FORM_sec_offset
    .byte    27                 # DW_AT_comp_dir
    .byte    14                 # DW_FORM_strp
    .byte    17                 # DW_AT_low_pc
    .byte    1                  # DW_FORM_addr
    .byte    18                 # DW_AT_high_pc
    .byte    6                  # DW_FORM_data4
    .byte    0                  # EOM(1)
    .byte    0                  # EOM(2)
    .byte    2                  # Abbreviation Code

```

```

.byte 46 # DW_TAG_subprogram
.byte 1 # DW_CHILDREN_yes
.byte 17 # DW_AT_low_pc
.byte 1 # DW_FORM_addr
.byte 18 # DW_AT_high_pc
.byte 6 # DW_FORM_data4
.byte 64 # DW_AT_frame_base
.byte 24 # DW_FORM_exprloc
.ascii "\227B" # DW_AT_GNU_all_call_sites
.byte 25 # DW_FORM_flag_present
.byte 3 # DW_AT_name
.byte 14 # DW_FORM_strp
.byte 58 # DW_AT_decl_file
.byte 11 # DW_FORM_data1
.byte 59 # DW_AT_decl_line
.byte 11 # DW_FORM_data1
.byte 39 # DW_AT_prototyped
.byte 25 # DW_FORM_flag_present
.byte 63 # DW_AT_external
.byte 25 # DW_FORM_flag_present
.byte 0 # EOM(1)
.byte 0 # EOM(2)
.byte 3 # Abbreviation Code
.byte 5 # DW_TAG_formal_parameter
.byte 0 # DW_CHILDREN_no
.byte 2 # DW_AT_location
.byte 23 # DW_FORM_sec_offset
.byte 3 # DW_AT_name
.byte 14 # DW_FORM_strp
.byte 58 # DW_AT_decl_file
.byte 11 # DW_FORM_data1
.byte 59 # DW_AT_decl_line
.byte 11 # DW_FORM_data1
.byte 73 # DW_AT_type
.byte 19 # DW_FORM_ref4
.byte 0 # EOM(1)
.byte 0 # EOM(2)
.byte 4 # Abbreviation Code
.byte 52 # DW_TAG_variable
.byte 0 # DW_CHILDREN_no
.byte 2 # DW_AT_location
.byte 23 # DW_FORM_sec_offset
.byte 3 # DW_AT_name
.byte 14 # DW_FORM_strp
.byte 58 # DW_AT_decl_file
.byte 11 # DW_FORM_data1
.byte 59 # DW_AT_decl_line
.byte 11 # DW_FORM_data1
.byte 73 # DW_AT_type
.byte 19 # DW_FORM_ref4
.byte 0 # EOM(1)
.byte 0 # EOM(2)
.byte 5 # Abbreviation Code
.byte 55 # DW_TAG_restrict_type
.byte 0 # DW_CHILDREN_no
.byte 73 # DW_AT_type
.byte 19 # DW_FORM_ref4
.byte 0 # EOM(1)

```

```

.byte 0 # EOM(2)
.byte 6 # Abbreviation Code
.byte 15 # DW_TAG_pointer_type
.byte 0 # DW_CHILDREN_no
.byte 73 # DW_AT_type
.byte 19 # DW_FORM_ref4
.byte 0 # EOM(1)
.byte 0 # EOM(2)
.byte 7 # Abbreviation Code
.byte 22 # DW_TAG_typedef
.byte 0 # DW_CHILDREN_no
.byte 73 # DW_AT_type
.byte 19 # DW_FORM_ref4
.byte 3 # DW_AT_name
.byte 14 # DW_FORM_strp
.byte 58 # DW_AT_decl_file
.byte 11 # DW_FORM_data1
.byte 59 # DW_AT_decl_line
.byte 11 # DW_FORM_data1
.byte 0 # EOM(1)
.byte 0 # EOM(2)
.byte 8 # Abbreviation Code
.byte 36 # DW_TAG_base_type
.byte 0 # DW_CHILDREN_no
.byte 3 # DW_AT_name
.byte 14 # DW_FORM_strp
.byte 62 # DW_AT_encoding
.byte 11 # DW_FORM_data1
.byte 11 # DW_AT_byte_size
.byte 11 # DW_FORM_data1
.byte 0 # EOM(1)
.byte 0 # EOM(2)
.byte 0 # EOM(3)
.section .debug_info,"",@progbits
.Lcu_begin0:
.long .Ldebug_info_end0-.Ldebug_info_start0 # Length of Unit
.Ldebug_info_start0:
.short 4 # DWARF version number
.long .debug_abbrev # Offset Into Abbrev. Section
.byte 8 # Address Size (in bytes)
.byte 1 # Abbrev [1] 0xb:0xa7 DW_TAG_compile_unit
.long .Linfo_string0 # DW_AT_producer
.short 12 # DW_AT_language
.long .Linfo_string1 # DW_AT_name
.long .Lline_table_start0 # DW_AT_stmt_list
.long .Linfo_string2 # DW_AT_comp_dir
.quad .Lfunc_begin0 # DW_AT_low_pc
.long .Lfunc_end0-.Lfunc_begin0 # DW_AT_high_pc
.byte 2 # Abbrev [2] 0x2a:0x43 DW_TAG_subprogram
.quad .Lfunc_begin0 # DW_AT_low_pc
.long .Lfunc_end0-.Lfunc_begin0 # DW_AT_high_pc
.byte 1 # DW_AT_frame_base
.byte 87 # DW_AT_GNU_all_call_sites

.long .Linfo_string3 # DW_AT_name
.byte 1 # DW_AT_decl_file
.byte 9 # DW_AT_decl_line
# DW_AT_prototyped

```

```

                                # DW_AT_external
        .byte    3                # Abbrev [3] 0x3f:0xf
DW_TAG_formal_parameter
        .long    .Ldebug_loc0     # DW_AT_location
        .long    .Linfo_string4   # DW_AT_name
        .byte    1                # DW_AT_decl_file
        .byte    9                # DW_AT_decl_line
        .long    109              # DW_AT_type
        .byte    3                # Abbrev [3] 0x4e:0xf
DW_TAG_formal_parameter
        .long    .Ldebug_loc1     # DW_AT_location
        .long    .Linfo_string8   # DW_AT_name
        .byte    1                # DW_AT_decl_file
        .byte    9                # DW_AT_decl_line
        .long    109              # DW_AT_type
        .byte    4                # Abbrev [4] 0x5d:0xf DW_TAG_variable
        .long    .Ldebug_loc2     # DW_AT_location
        .long    .Linfo_string9   # DW_AT_name
        .byte    1                # DW_AT_decl_file
        .byte    10               # DW_AT_decl_line
        .long    148              # DW_AT_type
        .byte    0                # End of Children Mark
        .byte    5                # Abbrev [5] 0x6d:0x5 DW_TAG_restrict_type
        .long    114              # DW_AT_type
        .byte    6                # Abbrev [6] 0x72:0x5 DW_TAG_pointer_type
        .long    119              # DW_AT_type
        .byte    7                # Abbrev [7] 0x77:0xb DW_TAG_typedef
        .long    130              # DW_AT_type
        .long    .Linfo_string7   # DW_AT_name
        .byte    3                # DW_AT_decl_file
        .byte    24               # DW_AT_decl_line
        .byte    7                # Abbrev [7] 0x82:0xb DW_TAG_typedef
        .long    141              # DW_AT_type
        .long    .Linfo_string6   # DW_AT_name
        .byte    2                # DW_AT_decl_file
        .byte    38               # DW_AT_decl_line
        .byte    8                # Abbrev [8] 0x8d:0x7 DW_TAG_base_type
        .long    .Linfo_string5   # DW_AT_name
        .byte    8                # DW_AT_encoding
        .byte    1                # DW_AT_byte_size
        .byte    7                # Abbrev [7] 0x94:0xb DW_TAG_typedef
        .long    159              # DW_AT_type
        .long    .Linfo_string12  # DW_AT_name
        .byte    3                # DW_AT_decl_file
        .byte    27               # DW_AT_decl_line
        .byte    7                # Abbrev [7] 0x9f:0xb DW_TAG_typedef
        .long    170              # DW_AT_type
        .long    .Linfo_string11  # DW_AT_name
        .byte    2                # DW_AT_decl_file
        .byte    45               # DW_AT_decl_line
        .byte    8                # Abbrev [8] 0xaa:0x7 DW_TAG_base_type
        .long    .Linfo_string10  # DW_AT_name
        .byte    7                # DW_AT_encoding
        .byte    8                # DW_AT_byte_size
        .byte    0                # End of Children Mark
.Ldebug_info_end0:
        .ident    "clang version 10.0.0-4ubuntu1 "
        .section   ".note.GNU-stack","",@progbits

```

```
.addrsig
.section .debug_line,"",@progbits
.Lline_table_start0:
```

内存未对齐。可以将所有元素移动一位来重新对齐b数组，或者使用偏移量。向量化后程序运行会更快。

**Write-up 5:** Check the assembly and verify that it does in fact vectorize properly. Also what do you notice when you run the command

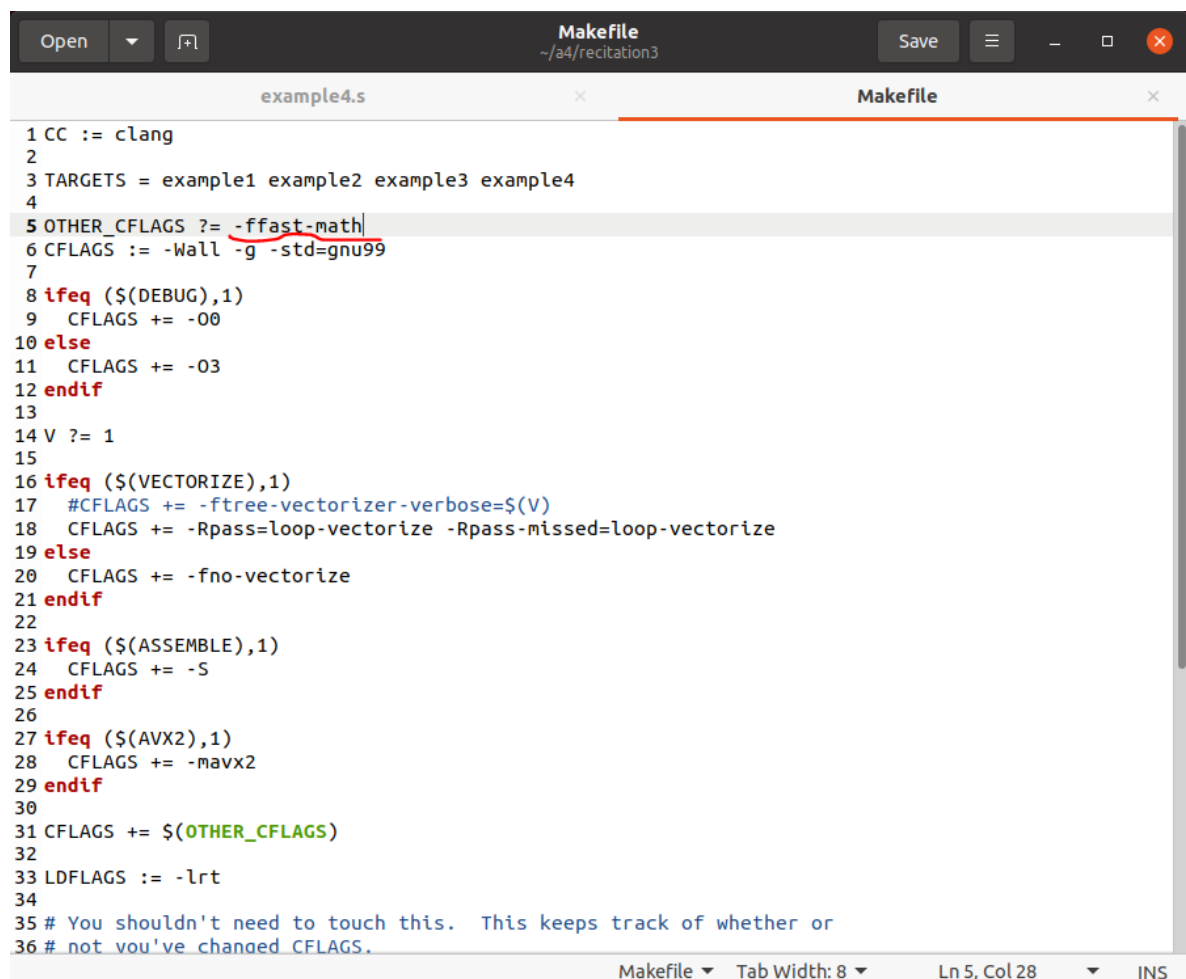
```
$ clang -O3 example4.c -o example4; ./example4
```

with and without the `-ffast-math` flag? Specifically, why do you see a difference in the output.

答：

对于汇编代码，材料已列举使用`addsd`命令的非向量化代码。

在makefile中添加`-ffast-math`



```
Makefile
~/a4/recitation3

example4.s
Makefile

1 CC := clang
2
3 TARGETS = example1 example2 example3 example4
4
5 OTHER_CFLAGS += -ffast-math
6 CFLAGS := -Wall -g -std=gnu99
7
8 ifeq ($(DEBUG),1)
9   CFLAGS += -O0
10 else
11   CFLAGS += -O3
12 endif
13
14 V ?= 1
15
16 ifeq ($(VECTORIZE),1)
17   #CFLAGS += -ftree-vectorizer-verbose=$(V)
18   CFLAGS += -Rpass=loop-vectorize -Rpass-missed=loop-vectorize
19 else
20   CFLAGS += -fno-vectorize
21 endif
22
23 ifeq ($(ASSEMBLE),1)
24   CFLAGS += -S
25 endif
26
27 ifeq ($(AVX2),1)
28   CFLAGS += -mavx2
29 endif
30
31 CFLAGS += $(OTHER_CFLAGS)
32
33 LDFLAGS := -lrt
34
35 # You shouldn't need to touch this. This keeps track of whether or
36 # not you've changed CFLAGS.
```

重新编译，可以看到确实应用了向量化：

```
zhangkeer@ubuntu: ~/a4/recitation3
zhangkeer@ubuntu:~/a4/recitation3$ make clean; make ASSEMBLE=1 VECTORIZE=1 example4.o
rm -rf
rm -f *.o *.s .cflags perf.data *.lst
clang -Wall -g -std=gnu99 -O3 -Rpass=loop-vectorize -Rpass-missed=loop-vectorize -S -ffast-math -c example4.c
example4.c:17:3: remark: vectorized loop (vectorization width: 2, interleaved count: 2) [-Rpass=loop-vectorize]
    for (i = 0; i < SIZE; i++) {
    ^
example4.c:25:3: remark: the cost-model indicates that vectorization is not beneficial [-Rpass-missed=loop-vectorize]
    for (int i = 0; i < SIZE; i++) {
    ^
example4.c:25:3: remark: the cost-model indicates that interleaving is not beneficial [-Rpass-missed=loop-vectorize]
example4.c:17:3: remark: vectorized loop (vectorization width: 2, interleaved count: 2) [-Rpass=loop-vectorize]
    for (i = 0; i < SIZE; i++) {
    ^
zhangkeer@ubuntu:~/a4/recitation3$
```

汇编代码中的命令也由addsd变成了addpd:

```
example4.s
~/a4/recitation3
1 .text
2 .file "example4.c"
3 .globl test # -- Begin function test
4 .p2align 4, 0x90
5 .type test,@function
6 test: # @test
7 .Lfunc_begin0:
8 .file 1 "/home/zhangkeer/a4/recitation3" "example4.c"
9 .loc 1 10 0 # example4.c:10:0
10 .cfi_startproc
11 # %bb.0:
12 #DEBUG_VALUE: test:a <- $rdi
13 #DEBUG_VALUE: test:a <- $rdi
14 xorpd %xmm0, %xmm0
15 xorl %eax, %eax
16 .Ltmp0:
17 #DEBUG_VALUE: test:i <- 0
18 #DEBUG_VALUE: test:y <- 0.000000e+00
19 #DEBUG_VALUE: test:x <- undef
20 xorpd %xmm1, %xmm1
21 .Ltmp1:
22 .p2align 4, 0x90
23 .LBB0_1: # =>This Inner Loop Header: Depth=1
24 #DEBUG_VALUE: test:y <- 0.000000e+00
25 #DEBUG_VALUE: test:i <- 0
26 #DEBUG_VALUE: test:a <- $rdi
27 .loc 1 18 7 prologue_end # example4.c:18:7
28 addpd (%rdi,%rax,8), %xmm0
29 addpd 16(%rdi,%rax,8), %xmm1
30 addpd 32(%rdi,%rax,8), %xmm0
31 addpd 48(%rdi,%rax,8), %xmm1
32 addpd 64(%rdi,%rax,8), %xmm0
33 addpd 80(%rdi,%rax,8), %xmm1
34 addpd 96(%rdi,%rax,8), %xmm0
35 addpd 112(%rdi,%rax,8), %xmm1
36 .Ltmp2:
37 .loc 1 17 26 # example4.c:17:26
```

不添加编译时标志 -ffast-math

```
clang -O3 example4.c -o example4; ./example4
```



输出结果为：

```
zhangkeer@ubuntu:~/a4/recitation3$ clang -O3 example4.c -o example4; ./example4
The decimal floating point sum result is: 11.667578
The raw floating point sum result is: 0x1.755cccec10aa5p+3
zhangkeer@ubuntu:~/a4/recitation3$
```

添加编译时标志 -ffast-math

```
clang -O3 -ffast-math example4.c -o example4; ./example4
```

输出结果为：

```
zhangkeer@ubuntu:~/a4/recitation3$ clang -O3 -ffast-math example4.c -o example4
; ./example4
The decimal floating point sum result is: 11.667578
The raw floating point sum result is: 0x1.755cccec10aa3p+3
zhangkeer@ubuntu:~/a4/recitation3$
```

输出结果不同，因为我们设置标志位允许了 clang 重新排序我们给它的操作。

**Write-up 6:** What speedup does the vectorized code achieve over the unvectorized code? What additional speedup does using `-mavx2` give? You may wish to run this experiment several times and take median elapsed times; you can report answers to the nearest 100% (e.g., 2×, 3×, etc). What can you infer about the bit width of the default vector registers on the `awsrun` machines? What about the bit width of the AVX2 vector registers? *Hint:* aside from speedup and the vectorization report, the most relevant information is that the data type for each array is `uint32_t`.

**Without vectorization:**

```
$ make
$ time ./loop
```

运行结果:

```
zhangkeer@ubuntu:~/a4/homework3$ make
clang -Wall -std=gnu99 -g -O3 -DNDEBUG -fno-vectorize -c loop.c
clang -o loop loop.o -lrt
zhangkeer@ubuntu:~/a4/homework3$ time ./loop
Elapsed execution time: 0.048396 sec; N: 1024, I: 100000, __OP__: +, __TYPE__: u
int32_t

real    0m0.050s
user    0m0.049s
sys     0m0.000s
zhangkeer@ubuntu:~/a4/homework3$
```

修改I的值，多次运行，使结果更有代表性：

```
zhangkeer@ubuntu:~/a4/homework3$ time ./loop
Elapsed execution time: 4.558956 sec; N: 1024, I: 1000000, __OP__: +, __TYPE__: uint32_t

real    0m4.560s
user    0m4.421s
sys     0m0.028s
zhangkeer@ubuntu:~/a4/homework3$
```

With vectorization :

```
$ make clean
$ make VECTORIZE=1
$ time ./loop
```

运行结果:

```
zhangkeer@ubuntu:~/a4/homework3$ make clean
rm -f loop *.o *.s .cflags perf.data */perf.data
zhangkeer@ubuntu:~/a4/homework3$ make VECTORIZE=1
clang -Wall -std=gnu99 -g -O3 -DNDEBUG -Rpass=loop-vectorize -Rpass-missed=loop-
vectorize -ffast-math -c loop.c
loop.c:70:9: remark: vectorized loop (vectorization width: 4, interleaved count:
  2) [-Rpass=loop-vectorize]
      for (j = 0; j < N; j++) {
      ^
clang -o loop loop.o -lrt
zhangkeer@ubuntu:~/a4/homework3$ time ./loop
Elapsed execution time: 0.021711 sec; N: 1024, I: 100000, __OP__: +, __TYPE__: u
int32_t

real    0m0.024s
user    0m0.024s
sys      0m0.000s
zhangkeer@ubuntu:~/a4/homework3$
```

修改I的值，多次运行，使结果更有代表性:

```
zhangkeer@ubuntu:~/a4/homework3$ time ./loop
Elapsed execution time: 1.133115 sec; N: 1024, I: 10000000, __OP__: +, __TYPE__: uint32_t

real    0m1.135s
user    0m1.134s
sys      0m0.000s
zhangkeer@ubuntu:~/a4/homework3$
```

With AVX instructions:

```
$ make clean
$ make VECTORIZE=1 AVX2=1
$ time ./loop
```

运行结果:

```
zhangkeer@ubuntu:~/a4/homework3$ make clean
rm -f loop *.o *.s .cflags perf.data */perf.data
zhangkeer@ubuntu:~/a4/homework3$ make VECTORIZE=1 AVX2=1
clang -Wall -std=gnu99 -g -O3 -DNDEBUG -Rpass=loop-vectorize -Rpass-missed=loop-
vectorize -ffast-math -mavx2 -c loop.c
loop.c:70:9: remark: vectorized loop (vectorization width: 8, interleaved count:
  4) [-Rpass=loop-vectorize]
      for (j = 0; j < N; j++) {
      ^
clang -o loop loop.o -lrt
zhangkeer@ubuntu:~/a4/homework3$ time ./loop
Elapsed execution time: 0.006121 sec; N: 1024, I: 100000, __OP__: +, __TYPE__: u
int32_t

real    0m0.008s
user    0m0.004s
sys      0m0.004s
zhangkeer@ubuntu:~/a4/homework3$
```

修改I的值，多次运行，使结果更有代表性：

```
zhangkeer@ubuntu:~/a4/homework3$ time ./loop
Elapsed execution time: 0.597142 sec; N: 1024, I: 10000000, __OP__: +, __TYPE__: uint32_t

real    0m0.599s
user    0m0.594s
sys     0m0.005s
zhangkeer@ubuntu:~/a4/homework3$
```

执行lscpu的结果是：

```
zhangkeer@ubuntu:~/a4/homework3$ lscpu
Architecture: x86_64
CPU op-mode(s): 32-bit, 64-bit
Byte Order: Little Endian
Address sizes: 45 bits physical, 48 bits virtual
CPU(s): 2
On-line CPU(s) list: 0,1
Thread(s) per core: 1
Core(s) per socket: 1
Socket(s): 2
NUMA node(s): 1
Vendor ID: GenuineIntel
CPU family: 6
Model: 142
Model name: Intel(R) Core(TM) i5-8250U CPU @ 1.60GHz
Stepping: 11
CPU MHz: 1799.997
BogoMIPS: 3599.99
Hypervisor vendor: VMware
Virtualization type: full
L1d cache: 64 KiB
L1i cache: 64 KiB
L2 cache: 512 KiB
L3 cache: 12 MiB
NUMA node0 CPU(s): 0,1
Vulnerability Itlb multihit: KVM: Mitigation: VMX unsupported
Vulnerability L1tr: Not affected
Vulnerability Mds: Vulnerable: Clear CPU buffers attempted, no mic
rocode; SMT Host state unknown
Vulnerability Meltdown: Not affected
Vulnerability Spec store bypass: Mitigation: Speculative Store Bypass disabled v
ia prctl and seccomp
Vulnerability Spectre v1: Mitigation: usercopy/swapgs barriers and __user
pointer sanitization
Vulnerability Spectre v2: Mitigation: Full generic retpoline, IBPB condit
ional, IBRS_FW, STIBP disabled, RSB filling
Vulnerability Srbds: Unknown: Dependent on hypervisor status
Vulnerability Tsx async abort: Not affected
Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtr
r pge mca cmov pat pse36 clflush mmx fxsr sse s
se2 ss syscall nx pdpe1gb rdtscp lm constant ts
c arch perfmon nopl xtopology tsc_reliable nonst
op_tsc cpuid pni pclmulqdq sse3 fma cx16 pcid
sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline
_timer aes xsave avx f16c rdrand hypervisor lah
f_lm abm 3dnowprefetch cpuid_fault invpcid sing
le ssbd ibrs ibpb stibp fsgsbase tsc_adjust bmi
1 avx2 smep bmi2 invpcid rdseed adx snap clflush
hopt xsaveopt xsavec xgetbv1 xsaves arat flush_
```

综合上面的结果，向量化代码将非向量化代码的运行速度提升了约四倍，使用AVX2向量寄存器又在向量化代码的基础上加速了两倍左右。

由于每个数组元素的数据类型是uint32\_t，默认向量寄存器的位宽应该是 $4 \times 32 = 128$ （性能提升四倍），AVX2 向量寄存器的位宽应该是 $2 \times 128 = 256$ （性能又提升两倍）。